

7avpxvvuh

February 17, 2024

IMPORTING LIBRARIES

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import sklearn
import seaborn as sns
from scipy import stats
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import KFold
from sklearn.linear_model import SGDRegressor
from sklearn.metrics import r2_score
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
from sklearn.linear_model import ElasticNet
from sklearn.preprocessing import PolynomialFeatures
```

LOADING DATASET

```
[6]: df = pd.read_csv(r'/winequality-red.csv', sep=",")
df
```

```
[6]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.4	0.700	0.00	1.9	0.076	
1	7.8	0.880	0.00	2.6	0.098	
2	7.8	0.760	0.04	2.3	0.092	
3	11.2	0.280	0.56	1.9	0.075	
4	7.4	0.700	0.00	1.9	0.076	
...	
1594	6.2	0.600	0.08	2.0	0.090	
1595	5.9	0.550	0.10	2.2	0.062	
1596	6.3	0.510	0.13	2.3	0.076	
1597	5.9	0.645	0.12	2.0	0.075	
1598	6.0	0.310	0.47	3.6	0.067	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	11.0	34.0	0.99780	3.51	0.56	

1	25.0	67.0	0.99680	3.20	0.68
2	15.0	54.0	0.99700	3.26	0.65
3	17.0	60.0	0.99800	3.16	0.58
4	11.0	34.0	0.99780	3.51	0.56
...
1594	32.0	44.0	0.99490	3.45	0.58
1595	39.0	51.0	0.99512	3.52	0.76
1596	29.0	40.0	0.99574	3.42	0.75
1597	32.0	44.0	0.99547	3.57	0.71
1598	18.0	42.0	0.99549	3.39	0.66

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5
...
1594	10.5	5
1595	11.2	6
1596	11.0	6
1597	10.2	5
1598	11.0	6

[1599 rows x 12 columns]

SUMMARIZING THE DATA

```
[7]: df.info() #no na values present
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
3   residual sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free sulfur dioxide    1599 non-null   float64
6   total sulfur dioxide   1599 non-null   float64
7   density                1599 non-null   float64
8   pH                    1599 non-null   float64
9   sulphates              1599 non-null   float64
10  alcohol                1599 non-null   float64
11  quality                1599 non-null   int64
dtypes: float64(11), int64(1)
```

memory usage: 150.0 KB

```
[8]: df.describe()
```

```
[8]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	\
count	1599.000000	1599.000000	1599.000000	1599.000000	
mean	8.319637	0.527821	0.270976	2.538806	
std	1.741096	0.179060	0.194801	1.409928	
min	4.600000	0.120000	0.000000	0.900000	
25%	7.100000	0.390000	0.090000	1.900000	
50%	7.900000	0.520000	0.260000	2.200000	
75%	9.200000	0.640000	0.420000	2.600000	
max	15.900000	1.580000	1.000000	15.500000	

	chlorides	free sulfur dioxide	total sulfur dioxide	density	\
count	1599.000000	1599.000000	1599.000000	1599.000000	
mean	0.087467	15.874922	46.467792	0.996747	
std	0.047065	10.460157	32.895324	0.001887	
min	0.012000	1.000000	6.000000	0.990070	
25%	0.070000	7.000000	22.000000	0.995600	
50%	0.079000	14.000000	38.000000	0.996750	
75%	0.090000	21.000000	62.000000	0.997835	
max	0.611000	72.000000	289.000000	1.003690	

	pH	sulphates	alcohol	quality
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	3.311113	0.658149	10.422983	5.636023
std	0.154386	0.169507	1.065668	0.807569
min	2.740000	0.330000	8.400000	3.000000
25%	3.210000	0.550000	9.500000	5.000000
50%	3.310000	0.620000	10.200000	6.000000
75%	3.400000	0.730000	11.100000	6.000000
max	4.010000	2.000000	14.900000	8.000000

PART A The summary of attributes is shown above,It can be observed that 1.there are no null values in the given dataset 2.All the values are continuous (float) and quality which is our target variable is int (discrete) 3.No categorical values are present in the dataset

```
[9]: wine_df=df.drop(['quality'],axis=1)
```

```
[10]: y=df['quality']
```

```
[11]: wine_df
```

```
[11]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.4	0.700	0.00	1.9	0.076	
1	7.8	0.880	0.00	2.6	0.098	
2	7.8	0.760	0.04	2.3	0.092	

3	11.2	0.280	0.56	1.9	0.075
4	7.4	0.700	0.00	1.9	0.076
...
1594	6.2	0.600	0.08	2.0	0.090
1595	5.9	0.550	0.10	2.2	0.062
1596	6.3	0.510	0.13	2.3	0.076
1597	5.9	0.645	0.12	2.0	0.075
1598	6.0	0.310	0.47	3.6	0.067

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates \
0	11.0	34.0	0.99780	3.51	0.56
1	25.0	67.0	0.99680	3.20	0.68
2	15.0	54.0	0.99700	3.26	0.65
3	17.0	60.0	0.99800	3.16	0.58
4	11.0	34.0	0.99780	3.51	0.56
...
1594	32.0	44.0	0.99490	3.45	0.58
1595	39.0	51.0	0.99512	3.52	0.76
1596	29.0	40.0	0.99574	3.42	0.75
1597	32.0	44.0	0.99547	3.57	0.71
1598	18.0	42.0	0.99549	3.39	0.66

	alcohol
0	9.4
1	9.8
2	9.8
3	9.8
4	9.4
...	...
1594	10.5
1595	11.2
1596	11.0
1597	10.2
1598	11.0

[1599 rows x 11 columns]

```
[12]: y.value_counts() #counting unique values of target variable shows imbalance
```

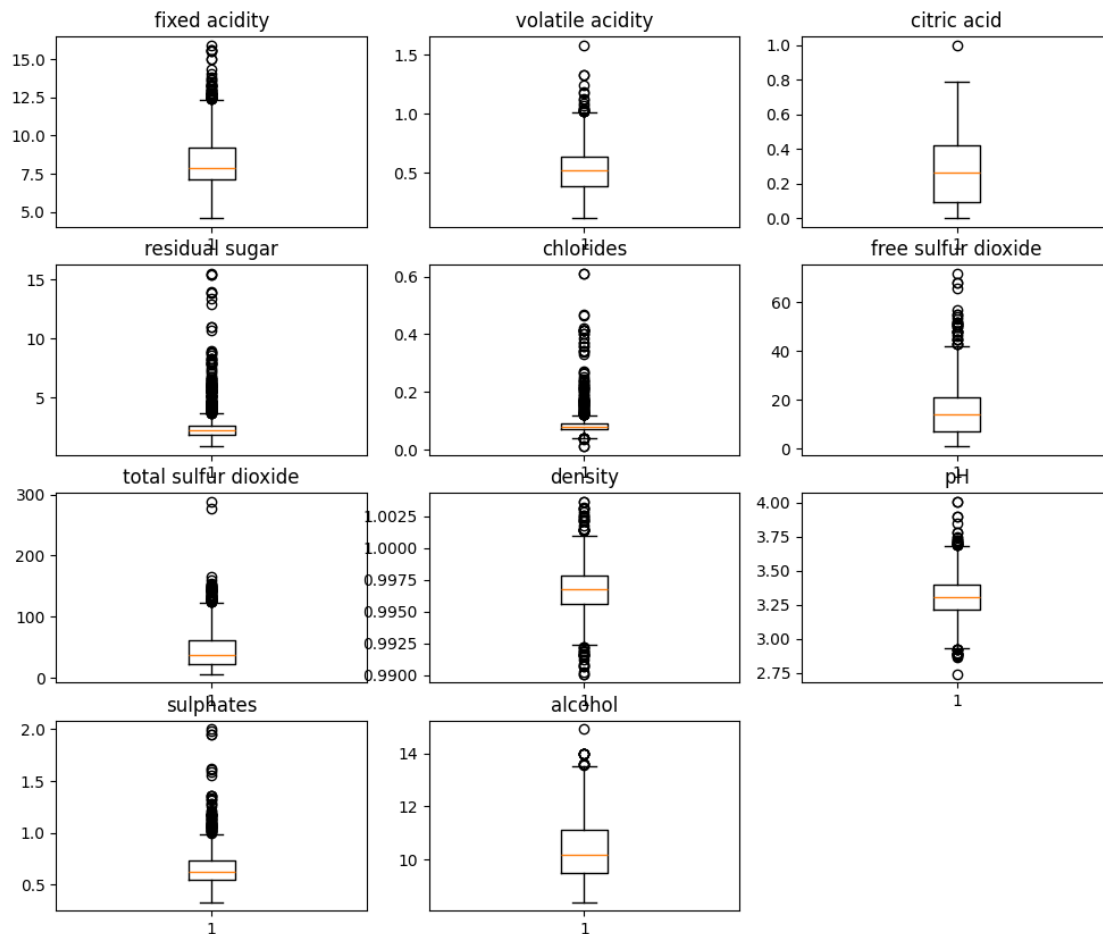
```
[12]: quality
5    681
6    638
7    199
4     53
8     18
3     10
Name: count, dtype: int64
```

```
[13]: #applying boxcox transformation to the target variable for normalizing
      ↪ distribution
```

```
from scipy.stats import boxcox
y_transformed, lambda_value = boxcox(y)
y1 = pd.Series(y_transformed)
```

```
[14]: #boxplot shows the outliers and data distribution of each attribute
```


```
%matplotlib inline
plt.figure(figsize=(12,10))
for i,c in enumerate(wine_df.columns):
    if i<=11:
        plt.subplot(4,3,i+1)
        plt.title(c)
        plt.boxplot(wine_df[c])
plt.show()
```



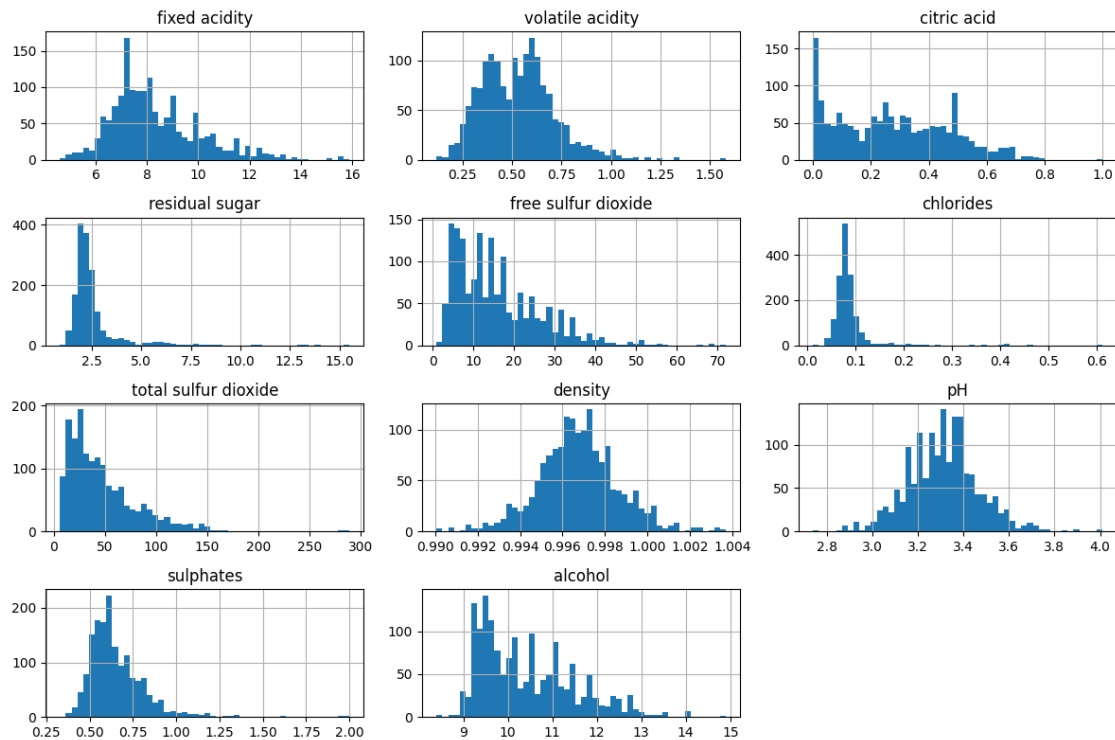
HISTOGRAM

```
[15]: df['quality'].hist(bins=30)
```

```
[15]: <Axes: >
```



```
[16]: wine_df[['fixed acidity', 'volatile acidity', 'citric acid', 'residual_
↳ sugar', 'free sulfur dioxide', 'chlorides', 'total sulfur_
↳ dioxide', 'density', 'pH', 'sulphates', 'alcohol']].hist(bins=50, figsize=(12,
↳ 8), layout=(4, 3))
plt.tight_layout()
plt.show()
```



PART B Histogram of various attributes including the target variable depicting the distribution of data. From the figure, most of the data is uniformly distributed; the target variable has an uneven distribution. Since all the attributes are continuous-valued with no null values, they do not require any special treatment. Outliers were not removed to prevent data loss, and they tell they provide the true nature of the dataset.

SCATTER PLOT

```
[ ]: sns.pairplot(df, kind='scatter')

# Show the plot
plt.show()
```

```
[ ]: corr = df.corr(method='pearson', numeric_only=True)
f, ax = plt.subplots(figsize=(9, 8))
sns.heatmap(corr, mask=np.zeros_like(corr, dtype=bool),
            cmap=sns.diverging_palette(220, 10, as_cmap=True),
            square=True, ax=ax)
```

```
[ ]: corr
```

```
[ ]: corr["quality"].sort_values(ascending=False) #correlation of target variable
      ↪ with the attributes
```

```
[ ]: #wine_df=wine_df.drop(['residual sugar'],axis=1)
```

PART C scatter plots showing relationship between various attributes which are generated using pearson coefficients. The above table data shows that residual sugar is least correlated with the target variable i.e 0.01 while attributes such as alcohol and volatile acidity contributes the most towards determining wine quality

SPLITTING DATASET FOR TRAINING

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(wine_df, y1, test_size=0.2,
↳ random_state=42)
```

```
[ ]: X_test.describe()
```

```
[ ]: wine_df.describe()
```

PART D Comparing the statistical values of the dataset and the test set we can observe that the test set represents the entire dataset. The dataset is splitted using random sampling which ensures that the random split is reproducible and that the test set is representative of the entire dataset

PART E PART 1 TRAINING MODEL USING NORMAL EQUATION AND K FOLD CROSS VALIDATION

```
[ ]: kf = KFold(shuffle=True,n_splits=4,random_state=42)
rmse_values = []
r2_values = []

for train_index, test_index in kf.split(wine_df):
    X_train, X_test = wine_df.iloc[train_index], wine_df.iloc[test_index]
    y_train, y_test = y1.iloc[train_index], y1.iloc[test_index]

    model = LinearRegression()
    model.fit(X_train, y_train)

    y_pred = model.predict(X_test)

    rmse = np.sqrt(mean_squared_error(y_test, y_pred))
    rmse_values.append(rmse)

    r2 = r2_score(y_test, y_pred)
    r2_values.append(r2)
    average_rmse = np.mean(rmse_values)
average_r2 = np.mean(r2_values)

print("Average RMSE:", average_rmse)
print("Average R^2:", average_r2)
```

SGD REGRESSOR


```
[ ]: model_sgd = SGDRegressor(max_iter=100, tol=1e-3, penalty='l2', alpha = 0.1)
tloss=[]
vloss=[]

for i in range(100):
    model_sgd.partial_fit(X_train, y_train)
    tloss.append(mean_squared_error(y_train, model_sgd.predict(X_train)))
    vloss.append(mean_squared_error(y_test, model_sgd.predict(X_test)))

plt.figure(figsize=(10,6))
plt.plot(tloss, label='Training Loss', marker='o')
plt.plot(vloss, label='Validation Loss')
plt.legend()
plt.show()
```

From the graph above it is evident that the model overfits the data depicted by the high validation loss and very less training loss.

```
[ ]: #DROPPING ONE FEATURE TO CHECK IF IT IMPROVES THE MODEL PERFORMANCE AND
    ↳RESIDUAL SUGAR HAS LEAST CORRELATION COEFFICIENT
wine_df1 = wine_df.drop(['residual sugar'],axis=1)
X_train, X_test, y_train, y_test = train_test_split(wine_df1, y, test_size=0.2,
    ↳random_state=42)
rmse_values = []
r2=[]
for train_index, test_index in kf.split(wine_df1):
    X_train, X_test = wine_df1.iloc[train_index], wine_df1.iloc[test_index]
    y_train, y_test = y1.iloc[train_index], y1.iloc[test_index]
    model = LinearRegression()
    model.fit(X_train, y_train)

    y_pred = model.predict(X_test)
    rmse = np.sqrt(mean_squared_error(y_test, y_pred))
    rmse_values.append(rmse)
    r2=r2_score(y_test, y_pred)

average_rmse = np.mean(rmse_values)
print(average_rmse)
r2
```

```
[ ]: model = SGDRegressor(max_iter=100, tol=1e-3, penalty='l2', alpha = 0.1)
tloss=[]
vloss=[]

for i in range(100):
    model.partial_fit(X_train, y_train)
    tloss.append(mean_squared_error(y_train, model.predict(X_train)))
```

```

        vloss.append(mean_squared_error(y_test, model.predict(X_test)))

plt.figure(figsize=(10,6))
plt.plot(tloss, label='Training Loss', marker='o')
plt.plot(vloss, label='Validation Loss')
plt.legend()
plt.show()

```

PART E RIDGE LASSO AND ELASTIC NET WITH DIFFERENT ALPHA VALUES

```
[ ]: alphas = [0.1, 1.0, 10.0] # Example alpha values
```

```
[ ]: # Train Ridge regression with different alpha values
for alpha in alphas:
    ridge = Ridge(alpha=alpha,max_iter=1000)
    ridge.fit(X_train, y_train)
    y_pred_ridge = ridge.predict(X_test)
    mse_ridge = mean_squared_error(y_test, y_pred_ridge)
    print(f"Mean Squared Error (Ridge) with alpha={alpha}: {mse_ridge}")
```

```
[ ]: for alpha in alphas:
    lasso = Lasso(alpha=alpha)
    lasso.fit(X_train, y_train)
    y_pred_lasso = lasso.predict(X_test)
    mse_lasso = mean_squared_error(y_test, y_pred_lasso)
    print(f"Mean Squared Error (Lasso) with alpha={alpha}: {mse_lasso}")
```

```
[ ]: elastic_net = ElasticNetCV(cv=3) #evaluating elasticnet using 4-fold cross_
    ↪validation
elastic_net.fit(X_train, y_train)
y_pred_elastic_net = elastic_net.predict(X_test)
mse_elastic_net = mean_squared_error(y_test, y_pred_elastic_net)
print("Mean Squared Error (Elastic Net):", mse_elastic_net)
```

PART E the lasso,ridge and elastic net regularization with different alpha values are given. RIDGE—>These MSE values represent the model's performance in terms of prediction accuracy for the respective alpha values. In this case, ridge with alpha=0.1 has the lowest MSE, suggesting that it provides the most accurate predictions among the three cases of alpha. Lasso with the learning rate of 0.1 also provides the least MSE among all the other values of alpha. For the above dataset Elastic net or ridge is best for regularization

POLYNOMIAL REGRESSION

```
[ ]: degree = 4
k = 4

train_errors = []
val_errors = []
```

```

kf = KFold(n_splits=k, shuffle=True)

for train_idx, val_idx in kf.split(wine_df):

    X_train, X_val = wine_df.iloc[train_idx], wine_df.iloc[val_idx]
    y_train, y_val = y1.iloc[train_idx], y1.iloc[val_idx]

    poly = PolynomialFeatures(degree=degree)
    X_train_poly = poly.fit_transform(X_train)
    X_val_poly = poly.transform(X_val)

    model = LinearRegression()
    model.fit(X_train_poly, y_train)

    y_train_pred = model.predict(X_train_poly)
    y_val_pred = model.predict(X_val_poly)

    train_mse = mean_squared_error(y_train, y_train_pred)
    val_mse = mean_squared_error(y_val, y_val_pred)
    r2=r2_score(y_val, y_val_pred)

    train_errors.append(train_mse)
    val_errors.append(val_mse)

plt.plot(range(1, k+1), train_errors, label='Train')
plt.plot(range(1, k+1), val_errors, label='Validation')
plt.xlabel('Fold')
plt.ylabel('MSE')
plt.legend()
plt.show()

```

```

[ ]: model_sgd = SGDRegressor(max_iter=100, tol=1e-3, penalty='l2', alpha = 0.1)
tloss=[]
vloss=[]

for i in range(100):
    model_sgd.partial_fit(X_train, y_train)
    tloss.append(mean_squared_error(y_train, model_sgd.predict(X_train)))
    vloss.append(mean_squared_error(y_test, model_sgd.predict(X_test)))

plt.figure(figsize=(10,6))
plt.plot(tloss, label='Training Loss', marker='o')
plt.plot(vloss, label='Validation Loss')
plt.legend()
plt.show()

```

The above polynomial model overfits the data which is depicted by the high validation loss and low training loss for both closed form and SGD method

RIDGE LASSO AND ELASTICNET FOR POLYNOMIAL REGRESSION

```
[ ]: for alpha in alphas:
    ridge = Ridge(alpha=alpha,max_iter=1000)
    ridge.fit(X_train, y_train)
    y_pred_ridge = ridge.predict(X_test)
    mse_ridge = mean_squared_error(y_test, y_pred_ridge)
    print(f"Mean Squared Error (Ridge) with alpha={alpha}: {mse_ridge}")
```

```
[ ]: for alpha in alphas:
    lasso = Lasso(alpha=alpha)
    lasso.fit(X_train, y_train)
    y_pred_lasso = lasso.predict(X_test)
    mse_lasso = mean_squared_error(y_test, y_pred_lasso)
    print(f"Mean Squared Error (Lasso) with alpha={alpha}: {mse_lasso}")
```

```
[ ]: elastic_net = ElasticNetCV(cv=3) #evaluating elasticnet using 4-fold cross
    validation
    elastic_net.fit(X_train, y_train)
    y_pred_elastic_net = elastic_net.predict(X_test)
    mse_elastic_net = mean_squared_error(y_test, y_pred_elastic_net)
    print("Mean Squared Error (Elastic Net):", mse_elastic_net)
```

the lasso,ridge and elastic net regularization with different alpha values are given. RIDGE→These MSE values represent the model's performance in terms of prediction accuracy for the respective alpha values. In this case, ridge with alpha=0.1 has the lowest MSE, suggesting that it provides the most accurate predictions among the three cases of alpha. Lasso with the learning rate of 0.1 also provides the least MSE among all the other values of alpha. For the above dataset Elastic net or ridge is best for regularization

PART G PREDICTION ON TEST LABELS

```
[ ]: ridge = Ridge(alpha=0.1,max_iter=1000)
    ridge.fit(X_train,y_train)
    y_pred_ridge = ridge.predict(X_test)
    mse_ridge = mean_squared_error(y_test, y_pred_ridge)
    r2=r2_score(y_test, y_pred_ridge)
    print("MSE VALUE:",mse_ridge)
    print("R2 VALUE:",r2)
```

The best model for regression on the above dataset is using the ridge regularizaion with alpha value=0.1 which reduces the MSE value to 0.1034389407556793.MSE is the best evaluation metric for the regression model above Future scope could include improving the R2 value of the model and increasing the dataset to prevent overfitting and evaluating more parameters using grid search or randomized search

[]:

[]:

[]:

[]: